

---

# NYC Taxi Trip Duration Project

---

## Final Project: Building Basic predictive models over the NYC Taxi Trip dataset.

You are provided with the NYC Taxi Trip Dataset. This dataset contains information about the taxi trips that took place in different parts of New York City and how much time did that trip take to complete.

In this project, the following are the tasks you must complete and submitted.

1. Choose the most suitable evaluation metric and state why you chose it.
2. Build a benchmark model for the given dataset.
3. Build a K-Nearest neighbours' model for the given dataset and find the best value of K.
4. Build a Linear model for the given dataset with regularisation. Attempt to interpret the variable coefficients of the Linear Model.
5. Build a Decision tree model for the given dataset. Attempt to interpret the variable importance.
6. Plot the following Bar plots:
  - A. train score of all the above models.
  - B. test (not validation!) score of all the above models.
  - C. Attempt to explain the observations from the plots (optional)

```
In [1]: #importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import warnings
warnings.filterwarnings("ignore")
sns.set()
```

### Reading the csv file

```
In [2]: data = pd.read_csv("nyc_taxi_trip_duration.csv")
```

### Checking the shape of dataset

```
In [3]: data.shape
```

```
Out[3]: (729322, 11)
```

### Checking first few rows of dataset

```
In [4]: data.head()
```

Out[4]:	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dr
0	id1080784	2	2016-02-29 16:40:21	2016-02-29 16:47:01	1	-73.953918	40.778873	
1	id0889885	1	2016-03-11 23:35:37	2016-03-11 23:53:57	2	-73.988312	40.731743	
2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48	2	-73.997314	40.721458	
3	id3744273	2	2016-01-05 09:44:31	2016-01-05 10:03:32	6	-73.961670	40.759720	
4	id0232939	1	2016-02-17 06:42:23	2016-02-17 06:56:31	1	-74.017120	40.708469	

### Checking if all id values are unique

```
In [5]: if data.id.nunique() == data.shape[0]:
        print("All id values are unique")
```

All id values are unique

### Checking for missing values

```
In [6]: data.isnull().sum()
```

```
Out[6]: id                                0
        vendor_id                        0
        pickup_datetime                  0
        dropoff_datetime                 0
        passenger_count                  0
        pickup_longitude                 0
        pickup_latitude                  0
        dropoff_longitude                0
        dropoff_latitude                 0
        store_and_fwd_flag               0
        trip_duration                    0
        dtype: int64
```

### Checking data types of the different features

```
In [7]: data.dtypes
```

```
Out[7]: id                                object
        vendor_id                        int64
        pickup_datetime                  object
        dropoff_datetime                 object
        passenger_count                  int64
        pickup_longitude                 float64
        pickup_latitude                  float64
        dropoff_longitude                float64
        dropoff_latitude                 float64
        store_and_fwd_flag               object
        trip_duration                    int64
        dtype: object
```

### Correcting the datatypes

```
In [8]: # pick-up datetime and dropoff_datetime should be datetime objects
        data['pickup_datetime'] = pd.to_datetime(data['pickup_datetime'])
        data['dropoff_datetime'] = pd.to_datetime(data['dropoff_datetime'])
```

```
In [9]: data.dtypes
```

```
Out[9]: id                                object
vendor_id                               int64
pickup_datetime                         datetime64[ns]
dropoff_datetime                        datetime64[ns]
passenger_count                         int64
pickup_longitude                        float64
pickup_latitude                         float64
dropoff_longitude                       float64
dropoff_latitude                        float64
store_and_fwd_flag                      object
trip_duration                           int64
dtype: object
```

### Checking if the trip\_duration matches with the difference of dropoff time and pickup time

```
In [10]: duration_diff = (data['trip_duration'] - ((data['dropoff_datetime'] - data['pickup_datetime']).total_seconds() / 60))
duration_diff
```

```
Out[10]: 0.0
```

- The trip\_duration is consistent with the pickup time and dropoff time.

### Creating pickup day of month, pickup day of week, and pickup hour features as these could be helpful for further analysis

```
In [11]: data['pickup_day'] = data.pickup_datetime.dt.day
data['pickup_dayow'] = data.pickup_datetime.dt.strftime('%A')
data['pickup_dow'] = data.pickup_datetime.dt.dayofweek
```

```
In [12]: data['pickup_hour'] = data.pickup_datetime.dt.hour
```

```
In [13]: data.head()
```

```
Out[13]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	id1080784	2	2016-02-29 16:40:21	2016-02-29 16:47:01	1	-73.953918	40.778873	-73.953918	40.778873
1	id0889885	1	2016-03-11 23:35:37	2016-03-11 23:53:57	2	-73.988312	40.731743	-73.988312	40.731743
2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48	2	-73.997314	40.721458	-73.997314	40.721458
3	id3744273	2	2016-01-05 09:44:31	2016-01-05 10:03:32	6	-73.961670	40.759720	-73.961670	40.759720
4	id0232939	1	2016-02-17 06:42:23	2016-02-17 06:56:31	1	-74.017120	40.708469	-74.017120	40.708469

### Shuffling and creating train and test data set

```
In [14]: from sklearn.utils import shuffle

# Shuffling the Dataset
data = shuffle(data, random_state = 42)
```

```
#creating 4 divisions
div = int(data.shape[0]/4)

# 3 parts to train set and 1 part to test set
train = data[: (3*div+1)]
test = data[3*div+1:]
```

In [15]: `train.shape`

Out[15]: (546991, 15)

In [16]: `test.shape`

Out[16]: (182331, 15)

## EDA

The EDA will be carried out on the training data set only to avoid any data leakage to the test dataset.

### Univariate Analysis

#### Getting an overview of the continuous data

In [17]: `train.describe()`

Out[17]:

	vendor_id	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	trip_d
<b>count</b>	546991.000000	546991.000000	546991.000000	546991.000000	546991.000000	546991.000000	5.469
<b>mean</b>	1.535572	1.664185	-73.973443	40.750949	-73.973383	40.751765	9.506
<b>std</b>	0.498733	1.314420	0.042147	0.034114	0.041840	0.035733	4.055
<b>min</b>	1.000000	0.000000	-79.569733	34.712234	-80.355431	32.181141	1.000
<b>25%</b>	1.000000	1.000000	-73.991852	40.737339	-73.991341	40.735851	3.970
<b>50%</b>	2.000000	1.000000	-73.981720	40.754086	-73.979767	40.754517	6.630
<b>75%</b>	2.000000	2.000000	-73.967346	40.768326	-73.963028	40.769741	1.076
<b>max</b>	2.000000	6.000000	-65.897385	51.881084	-65.897385	43.911762	1.939


- The trip duration has a very wide range and thus might contain outliers. Needs to be checked further.

In [18]:

```
# setting image resolution
plt.figure(figsize = (8,2) , dpi = 140)


# Plotting histogram and descriptive summary
sns.kdeplot(train['trip_duration'], shade = True)
```

Out[18]: <AxesSubplot:xlabel='trip\_duration', ylabel='Density'>



**Because of the presense of outliers, log of trip\_duration would be helpful in understanding the distribution.**

```
In [19]: log_trip_duration = np.log(train['trip_duration'] + 1)
plt.hist(log_trip_duration.values, bins=100)
plt.xlabel('log(trip_duration)')
plt.ylabel('number of train records')
plt.show()
```




```
In [20]: (train['trip_duration']/3600).describe() # Trip duration in hours
```

```
Out[20]: count      546991.000000
mean         0.264058
std          1.126499
min          0.000278
25%          0.110278
50%          0.184167
75%          0.298889
max          538.815556
Name: trip_duration, dtype: float64
```

```
In [21]: fig, ax = plt.subplots(figsize = (15,4), ncols=1, nrows=1)
ax.set_xlim(0,48)
sns.distplot(train['trip_duration']/3600,ax=ax,bins=1000,kde=False, hist_kws={'log':True})
```

```
Out[21]: <AxesSubplot:xlabel='trip_duration'>
```




- We can see that most of the trips are within the 24 hour duration, and the no. of trips increase with duration of about 24 hours, maybe because the meter resets at midnight, and maybe the drivers forget to reset their meters.
- We will restrict the trip duration to less than 22 hours for the reason stated above.
- There are a large number of trips with duration less than an hour. It would be good to have a deeper look at it.

```
In [22]: train = train[train['trip_duration'] < (3600*22)]
```

```
In [23]: fig, ax = plt.subplots(figsize = (15,4), ncols=1, nrows=1)
hour_data = train[train['trip_duration'] < 3600]
sns.distplot(hour_data['trip_duration']/60,ax=ax,bins=60,kde=False, hist_kws={'log':True},
```

```
Out[23]: <AxesSubplot:xlabel='trip_duration'>
```



- We can see that there are a number of trips with duration less than 60 seconds, which is highly improbable in real life. So, we'll remove those trips with trip\_duration less than 60.

```
In [24]: (train['trip_duration'] < 60).value_counts()
```

```
Out[24]: False    543102
         True      3210
         Name: trip_duration, dtype: int64
```

```
In [25]: train = train[train['trip_duration'] > 60]
```

```
In [26]: train.shape
```

Out[26]: (543027, 15)

### Checking the range for pickup and dropoff times

```
In [27]: train['pickup_datetime'].min()
```

Out[27]: Timestamp('2016-01-01 00:01:14')

```
In [28]: train['pickup_datetime'].max()
```

Out[28]: Timestamp('2016-06-30 23:59:37')

```
In [29]: train['dropoff_datetime'].min()
```

Out[29]: Timestamp('2016-01-01 00:05:54')

```
In [30]: train['dropoff_datetime'].max()
```

Out[30]: Timestamp('2016-07-01 00:46:37')

- The data is for a period of six months. The day of the month, and the day of the week might provide us some insights to predict the trip duration.

### Checking values of store and fwd flag

```
In [31]: train.store_and_fwd_flag.value_counts()
```

Out[31]:

N	540050
Y	2977

Name: store\_and\_fwd\_flag, dtype: int64

```
In [32]: # This feature could provide some insights at the later stage, so we can convert it into a binary feature  
train['store_and_fwd_flag'] = 1 * (train.store_and_fwd_flag.values == 'Y')
```

```
In [33]: train.store_and_fwd_flag.value_counts(normalize = True)
```

Out[33]:

0	0.994518
1	0.005482

Name: store\_and\_fwd\_flag, dtype: float64

```
In [34]: # Analysing Vendor Id  
train['vendor_id'].value_counts(normalize=True)
```

Out[34]:

2	0.535692
1	0.464308

Name: vendor\_id, dtype: float64

```
In [35]: # Analysing passenger counts  
train['passenger_count'].value_counts()
```

Out[35]:

1	384683
2	78452
5	29063
3	22228

```

6      18060
4      10532
0         9
Name: passenger_count, dtype: int64

```

In [36]:

```

# Plotting vendor id, passenger counts, and store and fwd flag


fig, (ax1,ax2,ax3) = plt.subplots(1,3, figsize = (15,6))

sns.countplot(x="vendor_id", data=train, palette="Set1", ax = ax1, order = data['vendor_id'].value_counts().index)
sns.countplot(x="passenger_count", data=train, palette="Reds_r", ax = ax2, order = data['passenger_count'].value_counts().index)
sns.countplot(x="store_and_fwd_flag", data=train, palette="Set1", ax = ax3)

ax1.bar_label(container=ax1.containers[0])
ax2.bar_label(container=ax2.containers[0])
ax3.bar_label(container=ax3.containers[0])

fig.tight_layout()

```



- Vendor 2 has more number of trips compared to vendor 1.
- Most of the trips have 1 passengers.
- We can see that there are a few rides with the passenger\_count as zero. We need to remove these from our training set.

In [37]:

```
train = train[train['passenger_count'] != 0]
```

In [38]:

```
train.shape
```

Out[38]:

```
(543018, 15)
```

In [39]:

```

plt.figure(figsize=(22, 6))

# Plotting pickups by day of week
plt.subplot(121)
sns.countplot(train['pickup_dayow'], order = train['pickup_dayow'].value_counts().index, palette="Set1")
plt.xlabel('Week Day')
plt.ylabel('Total Number of pickups')


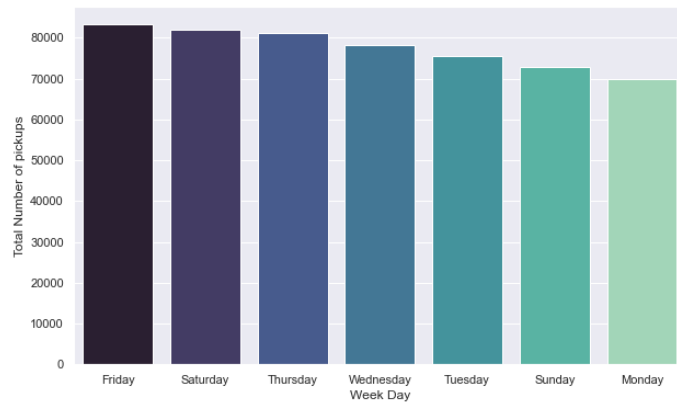
# Plotting pickups by hour of day
plt.subplot(122)
sns.countplot(train['pickup_hour'], color = 'b')

```



```
plt.xlabel('Hour of Day')
plt.ylabel('Total number of pickups')
```

Out[39]: Text(0, 0.5, 'Total number of pickups')



- We can see that the least number of pickups happen on Mondays while Fridays have the highest number of pickups.
- Early morning pickups are low, and pickups are more after 08:00 AM which can be considered as the start of work hours.

## Bivariate and multivariate analysis

In [40]:


```
grouped_data1 = train.groupby(['pickup_hour', 'vendor_id'])['trip_duration'].mean().reset_index()
grouped_data2 = train.groupby(['pickup_dow', 'vendor_id'])['trip_duration'].mean().reset_index()
```

In [41]:

```
fig, ax = plt.subplots(ncols=2, figsize = (12,6), sharey=True)
sns.lineplot(data = grouped_data1, x = 'pickup_hour', y = 'trip_duration', hue = 'vendor_id')
sns.lineplot(data = grouped_data2, x = 'pickup_dow', y = 'trip_duration', hue = 'vendor_id')

ax[0].set_xlabel('hour')
ax[1].set_xlabel('weekday')
ax[0].set_ylabel('trip duration')


fig.tight_layout()
```



- Trip durations are definitely shorter for late night and early morning hours that can be attributed to low traffic density
- It follows a similar pattern when compared to number of pickups indicating a correlation between number of pickups and trip duration
- Trip durations are particularly low on Saturday(5) and Sunday(6).
- Trip durations for vendor 2 are on an average higher than trip durations for vendor 1 on all days.

In [128...]

```
train.passenger_count.value_counts()
plt.figure(figsize=(22, 6))
train_sub = train[train['trip_duration'] < 3600]
sns.boxplot(x="passenger_count", y="trip_duration", data=train_sub, palette = 'Set2')
plt.show()
```



- The boxplot clearly shows that there is not much of a difference in distribution for the most frequently occurring passenger count values.

In [42]:

```
train[['pickup_longitude', 'pickup_latitude']].describe()
```

Out[42]:

	pickup_longitude	pickup_latitude
count	543018.000000	543018.000000
mean	-73.973567	40.750996

	<b>pickup_longitude</b>	<b>pickup_latitude</b>
<b>std</b>	0.041523	0.033848
<b>min</b>	-79.569733	34.712234
<b>25%</b>	-73.991859	40.737400
<b>50%</b>	-73.981735	40.754116
<b>75%</b>	-73.967438	40.768330
<b>max</b>	-65.897385	51.881084

- We have outliers in latitude and longitude as well.

In [43]:

```
xlong = train['pickup_longitude'].mean()
x_sd = train['pickup_longitude'].std()
ylat = train['pickup_latitude'].mean()
y_sd = train['pickup_latitude'].std()
```

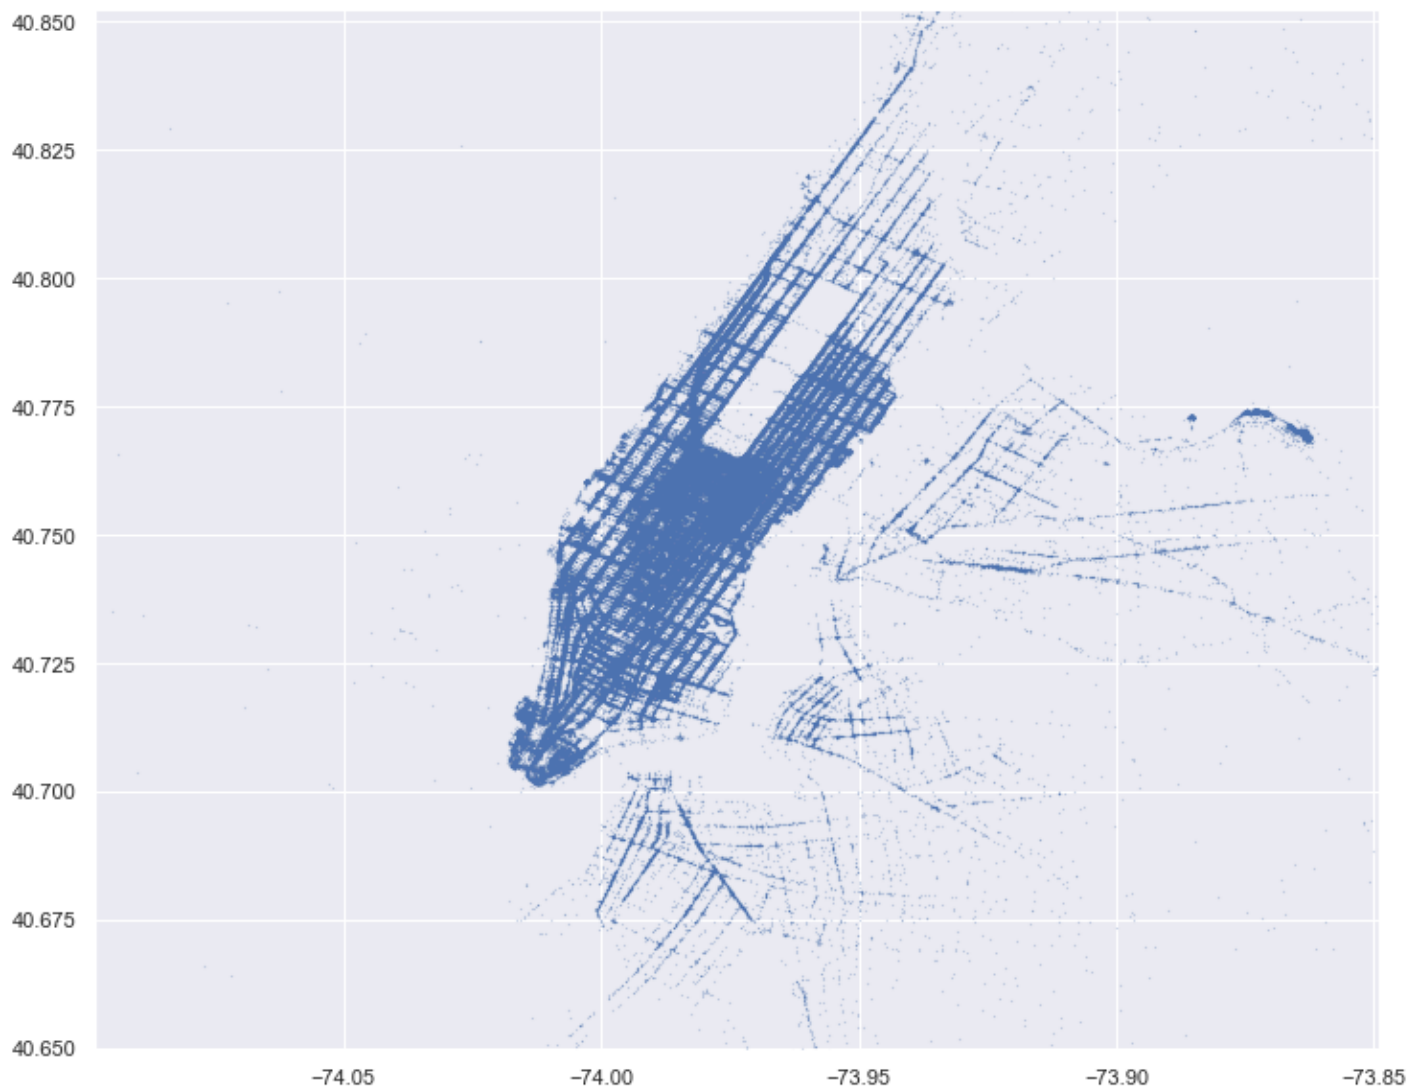
- We'll take points only within the +/- 3 sigma range to visualise the coordinates as this will eliminate the outliers.

In [44]:

```
fig, ax = plt.subplots(ncols=1, nrows=1, figsize=(12,10))
plt.ylim(ylat - 3*y_sd, ylat + 3*y_sd)
plt.xlim(xlong - 3*x_sd, xlong + 3*x_sd)
ax.scatter(train['pickup_longitude'], train['pickup_latitude'], s=0.01, alpha=1)
```

Out[44]:

```
<matplotlib.collections.PathCollection at 0x23d4560c370>
```



```
In [45]: train1 = train.loc[(train.pickup_latitude > 40.6) & (train.pickup_latitude < 40.9)]
train1 = train.loc[(train.pickup_longitude > -74.05) & (train.pickup_longitude < -73.7)]
test1 = test.loc[(test.pickup_latitude > 40.6) & (test.pickup_latitude < 40.9)]
test1 = test.loc[(test.pickup_longitude > -74.05) & (test.pickup_longitude < -73.7)]
```

## Clustering the pickup locations

```
In [46]: import geopandas as gpd
from shapely.geometry import Point, Polygon
```

```
In [47]: geometry = [Point(xy) for xy in zip(train1['pickup_longitude'], train1['pickup_latitude'])]
geometry[:3]
```

```
Out[47]: [<shapely.geometry.point.Point at 0x23d491c80d0>,
<shapely.geometry.point.Point at 0x23d4563bca0>,
<shapely.geometry.point.Point at 0x23d4563bcd0>]
```

```
In [48]: geo_df = gpd.GeoDataFrame(train1, crs = {'init': 'epsg:4326'}, geometry = geometry)
geo_df.head()
```

```
Out[48]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
469114	id2380741	2	2016-05-21 10:40:14	2016-05-21 10:51:11	1	-73.981796	40.76203

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
<b>694852</b>	id3946961	2	2016-01-08 18:49:27	2016-01-08 18:52:42	5	-73.980965	40.74767
<b>696324</b>	id0833913	1	2016-05-22 00:54:10	2016-05-22 01:08:10	1	-73.951065	40.78272
<b>356496</b>	id1336849	1	2016-06-11 10:32:12	2016-06-11 10:38:50	1	-73.987625	40.76279
<b>645318</b>	id1610858	1	2016-04-03 10:45:51	2016-04-03 10:57:13	3	-73.964333	40.79250

In [49]:

```
from sklearn.cluster import KMeans
k_means = KMeans(n_clusters = 50)
df_pick = train1[['pickup_longitude', 'pickup_latitude']]
df_drop = train1[['dropoff_longitude', 'dropoff_latitude']]

k_means.fit(df_pick)
```

Out[49]:

▼ KMeans  
KMeans(n\_clusters=50)

In [50]:

```
clust_pick = k_means.labels_
train1['label_pick'] = clust_pick.tolist()
train1['label_drop'] = k_means.predict(df_drop)

centroid_pickups = pd.DataFrame(k_means.cluster_centers_, columns = ['centroid_pick_long',
centroid_dropoff = pd.DataFrame(k_means.cluster_centers_, columns = ['centroid_drop_long',

centroid_pickups.shape
centroid_pickups['label_pick'] = centroid_pickups.index
centroid_dropoff['label_drop'] = centroid_dropoff.index

train1 = pd.merge(train1, centroid_pickups, how='left', on=['label_pick'])
train1 = pd.merge(train1, centroid_dropoff, how='left', on=['label_drop'])
train1.head()
```

Out[50]:

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dr
<b>0</b>	id2380741	2	2016-05-21 10:40:14	2016-05-21 10:51:11	1	-73.981796	40.762035	
<b>1</b>	id3946961	2	2016-01-08 18:49:27	2016-01-08 18:52:42	5	-73.980965	40.747677	
<b>2</b>	id0833913	1	2016-05-22 00:54:10	2016-05-22 01:08:10	1	-73.951065	40.782722	
<b>3</b>	id1336849	1	2016-06-11 10:32:12	2016-06-11 10:38:50	1	-73.987625	40.762791	

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_latitude
4	id1610858	1	2016-04-03 10:45:51	2016-04-03 10:57:13	3	-73.964333	40.792503	

5 rows × 22 columns

```
In [51]: summary_avg_time = pd.DataFrame(train1.groupby('label_pick')['trip_duration'].mean())
summary_avg_time.reset_index(inplace = True)
summary_pref_clus = pd.DataFrame(train1.groupby(['label_pick', 'label_drop'])['id'].count())
summary_pref_clus = summary_pref_clus.reset_index()
summary_pref_clus = summary_pref_clus.loc[summary_pref_clus.groupby('label_pick')['id'].count().sort_values(ascending=False).index]
summary = pd.merge(summary_avg_time, summary_pref_clus, how = 'left', on = 'label_pick')
summary = summary.rename(columns={'trip_duration': 'avg_trip_time'})
summary.head()
```

```
Out[51]:
```

	label_pick	avg_trip_time	label_drop	id
0	0	817.360420	22	2001
1	1	2594.393606	38	236
2	2	707.563587	14	1282
3	3	1930.848292	38	378
4	4	815.173536	12	1019

```
In [52]: # Summary of distribution of data over 50 clusters
summary_full_data = pd.DataFrame(train1.groupby('label_pick')['id'].count())
summary_full_data['id'].describe()
```

```
Out[52]: count      50.000000
mean      10858.040000
std       7658.432426
min         8.000000
25%       2995.000000
50%      12339.500000
75%      16545.250000
max      26180.000000
Name: id, dtype: float64
```

```
In [53]: def clusters_map(clus_data, full_data, tile = 'OpenStreetMap', sig = 0, zoom = 12, circle = 1):
    """ function to plot clusters on map """
    map_1 = folium.Map(location=[40.767937, -73.982155], zoom_start=zoom, tiles= tile) # 'OpenStreetMap'
    summary_full_data = pd.DataFrame(full_data.groupby('label_pick')['id'].count())
    summary_full_data.reset_index(inplace = True)
    if sig == 1:
        summary_full_data = summary_full_data.loc[summary_full_data['id']>11500]
        sig_cluster = summary_full_data['label_pick'].tolist()
        clus_summary = summary_full_data[summary_full_data['label_pick'].isin(sig_cluster)]

    for i in sig_cluster:
        pick_long = clus_data.loc[clus_data.index == i]['centroid_pick_long'].values[0]
        pick_lat = clus_data.loc[clus_data.index == i]['centroid_pick_lat'].values[0]
        clus_no = clus_data.loc[clus_data.index == i]['label_pick'].values[0]
        most_visited_clus = clus_summary.loc[clus_summary['label_pick'] == i]['label_drop'].values[0]
        avg_trip_time = clus_summary.loc[clus_summary['label_pick'] == i]['avg_trip_time'].values[0]
        pop = 'cluster = ' + str(clus_no) + ' & most visited cluster = ' + str(most_visited_clus)
        if circle == 1:
            folium.CircleMarker(location=[pick_lat, pick_long], radius=radius, fill_color='red', fill_opacity=0.5, stroke='red', stroke_weight=2)
```



```

        color='#F08080',
        fill_color='#3186cc', popup=pop).add_to(map_1)
    folium.Marker([pick_lat, pick_long], popup=pop).add_to(map_1)
    return map_1

```

```
In [54]: import folium
```

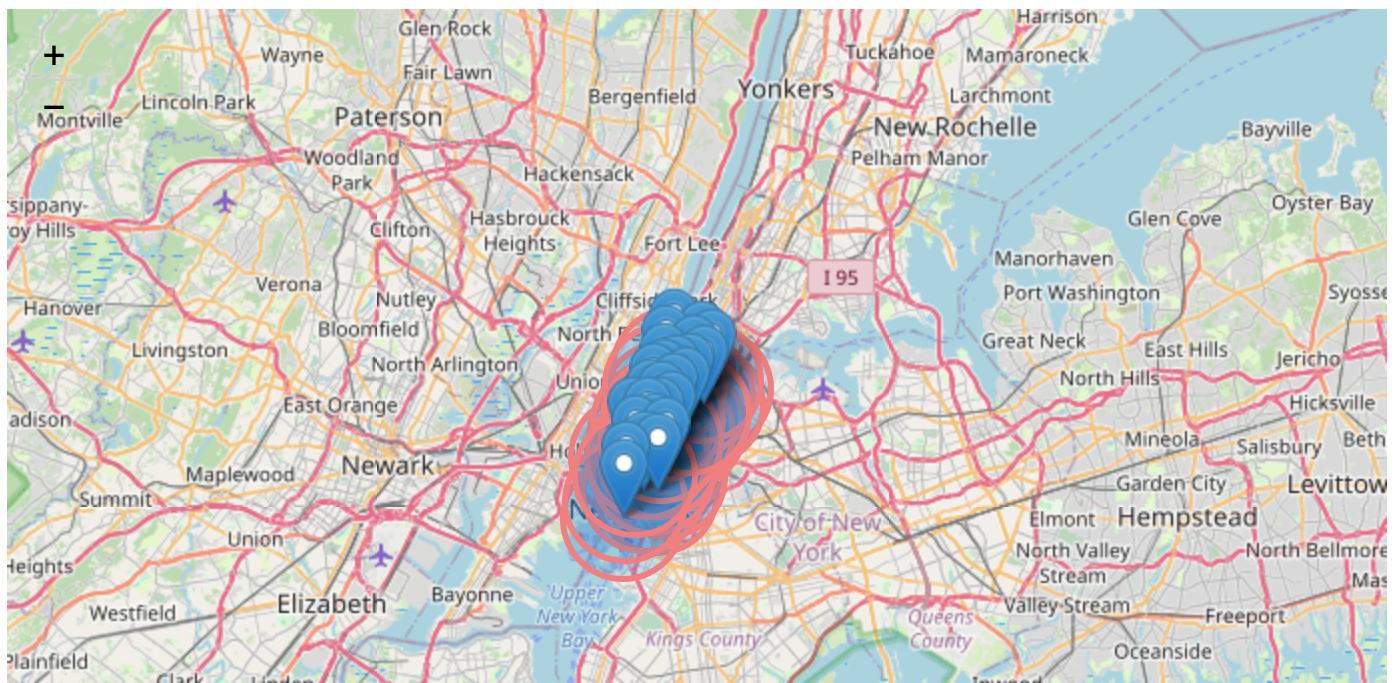
```
In [55]: clus_map = clusters_map(centroid_pickups, train1, sig =0, zoom =12, circle =1, tile = 'Stamen')
clus_map
```

Out[55]:



```
In [56]: # Visualizing top 50% of the clusters
clus_map_sig = clusters_map(centroid_pickups, train1, sig =1, circle =1, zoom = 12)
clus_map_sig
```

Out[56]:






In [57]:

```
# Checking correlation

# plotting heatmap usill all methods for all numerical variables
plt.figure(figsize=(36,6), dpi=140)
for j,i in enumerate(['pearson','kendall','spearman']):
    plt.subplot(1,3,j+1)
    correlation = train.dropna().corr(method=i)
    # Generate a mask for the upper triangle
    mask = np.triu(np.ones_like(correlation, dtype=bool))
    # Generate a custom diverging colormap
    cmap = sns.diverging_palette(220, 10, as_cmap=True)
    sns.heatmap(correlation, mask = mask, cmap = cmap, linewidth = 2)
    plt.title(i, fontsize=18)
```



- The correlation plot shows that the strongest correlations are between `passenger_count` and `vendor_id`, `dropoff_longitude` and `pickup_longitude`, `dropoff_latitude` and `pickup_latitude`.

## Feature Engineering

Knowing the distance between pickup and dropoff could be helpful in predicting the trip duration. This can be calculated using the haversine distance formula that uses the latitude and longitude coordinates to calculate the distance between the two places.

In [58]:

```
def haversine(lat1, lng1, lat2, lng2):
    lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
    AVG_EARTH_RADIUS = 6371 # in km
    lat = lat2 - lat1
    lng = lng2 - lng1
    d = np.sin(lat * 0.5) ** 2 + np.cos(lat1) * np.cos(lat2) * np.sin(lng * 0.5) ** 2
    h = 2 * AVG_EARTH_RADIUS * np.arcsin(np.sqrt(d))
    return h
```

In [59]:

```
train['distance'] = haversine(train.pickup_latitude, train.pickup_longitude,
                              train.dropoff_latitude, train.dropoff_longitude)
test['distance'] = haversine(test.pickup_latitude, test.pickup_longitude,
                             test.dropoff_latitude, test.dropoff_longitude)
```

In [60]:

```
train.distance.describe()
```




```
Out[60]: count      543018.000000
          mean         3.460875
          std         4.398021
          min         0.000000
          25%         1.247133
          50%         2.112514
          75%         3.901750
          max        1240.908677
          Name: distance, dtype: float64
```

```
In [61]: train.loc[:, 'avg_speed_h'] = train['distance'] / (train['trip_duration']/3600)
```

```
In [62]: fig, ax = plt.subplots(figsize = (12,4), ncols=2, sharey=True)
          ax[0].plot(train.groupby('pickup_hour').mean()['avg_speed_h'], 'bo-', lw=2, alpha=0.7)
          ax[1].plot(train.groupby('pickup_dow').mean()['avg_speed_h'], 'go-', lw=2, alpha=0.7)

          ax[0].set_xlabel('hour')
          ax[1].set_xlabel('weekday')
          ax[1].set_xticks([0,1,2,3,4,5,6], ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
          ax[0].set_ylabel('average speed (km/hr)')
          fig.suptitle('Average traffic speed')
          fig.tight_layout()
```



- The hours 08:00 to 18:00 seem to be the work hours with more traffic and thus lower traffic speed.
- A new feature named work\_hours can be introduced with its value as 1 during the work hours, from Monday to Friday, and 0 otherwise.

```
In [63]: train['work_hours'] = ((train['pickup_hour'] >= 8) & (train['pickup_hour'] < 18) & (train
train.head()
```

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
<b>469114</b>	id2380741	2	2016-05-21 10:40:14	2016-05-21 10:51:11	1	-73.981796	40.76203
<b>694852</b>	id3946961	2	2016-01-08 18:49:27	2016-01-08 18:52:42	5	-73.980965	40.74767
<b>696324</b>	id0833913	1	2016-05-22 00:54:10	2016-05-22 01:08:10	1	-73.951065	40.78272
<b>356496</b>	id1336849	1	2016-06-11 10:32:12	2016-06-11 10:38:50	1	-73.987625	40.76279
<b>645318</b>	id1610858	1	2016-04-03 10:45:51	2016-04-03 10:57:13	3	-73.964333	40.79250

# Evaluation metric

- The most suitable evaluation metric for trip duration would be RMSLE, as it wouldn't be affected by outliers.

```
In [64]: def rmsle(evaluator,X,real):
          sum = 0.0
          predicted = evaluator.predict(X)
          #print("Number predicted less than 0: {}".format(np.where(predicted < 0)[0].shape))

          predicted[predicted < 0] = 0
          for x in range(len(predicted)):
              p = np.log(predicted[x]+1)
              r = np.log(real[x]+1)
              sum = sum + (p-r)**2
          return (sum/len(predicted))**0.5
```

## Preparing Test Data

```
In [65]: test.head()
```

```
Out[65]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
<b>409888</b>	id3315205	1	2016-06-21 22:06:35	2016-06-21 22:19:19	1	-73.981903	40.76569
<b>666838</b>	id0273627	1	2016-01-29 08:50:01	2016-01-29 09:37:45	1	-73.963600	40.77439
<b>421168</b>	id3291472	1	2016-03-30 12:36:29	2016-03-30 12:45:53	1	-73.991081	40.73741
<b>348868</b>	id2444699	1	2016-04-11 05:36:57	2016-04-11 05:44:28	1	-73.962936	40.76647
<b>34687</b>	id2159293	1	2016-01-01 02:52:32	2016-01-01 03:01:51	1	-73.976067	40.75044

```
In [66]: test['store_and_fwd_flag'] = 1 * (test.store_and_fwd_flag.values == 'Y')
```

```
In [67]: test['work_hours'] = ((test['pickup_hour'] >= 8) & (test['pickup_hour'] < 18) & (test['pic
```

```
In [68]: test.loc[:, 'avg_speed_h'] = 1000 * test['distance'] / test['trip_duration']
```

```
In [69]: test.head()
```

```
Out[69]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
<b>409888</b>	id3315205	1	2016-06-21 22:06:35	2016-06-21 22:19:19	1	-73.981903	40.76569
<b>666838</b>	id0273627	1	2016-01-29 08:50:01	2016-01-29 09:37:45	1	-73.963600	40.77439

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
<b>421168</b>	id3291472	1	2016-03-30 12:36:29	2016-03-30 12:45:53	1	-73.991081	40.73741
<b>348868</b>	id2444699	1	2016-04-11 05:36:57	2016-04-11 05:44:28	1	-73.962936	40.76647
<b>34687</b>	id2159293	1	2016-01-01 02:52:32	2016-01-01 03:01:51	1	-73.976067	40.75044

## Benchmark Model

```
In [70]: # storing simple median in a new column in the test set as "simple_median"
test['median_duration'] = train['trip_duration'].median()
test.head()
```

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
<b>409888</b>	id3315205	1	2016-06-21 22:06:35	2016-06-21 22:19:19	1	-73.981903	40.76569
<b>666838</b>	id0273627	1	2016-01-29 08:50:01	2016-01-29 09:37:45	1	-73.963600	40.77439
<b>421168</b>	id3291472	1	2016-03-30 12:36:29	2016-03-30 12:45:53	1	-73.991081	40.73741
<b>348868</b>	id2444699	1	2016-04-11 05:36:57	2016-04-11 05:44:28	1	-73.962936	40.76647
<b>34687</b>	id2159293	1	2016-01-01 02:52:32	2016-01-01 03:01:51	1	-73.976067	40.75044

```
In [71]: #calculating mean absolute error
from sklearn.metrics import mean_squared_error as MSE

median_rmsle = MSE(np.log(test['trip_duration']+1) , np.log(test['median_duration']+1), sc
median_rmsle
```

Out[71]: 0.7965105610441814

## KNN Model

```
In [72]: train.dtypes
```

```
Out[72]: id                                object
vendor_id                             int64
pickup_datetime                       datetime64[ns]
dropoff_datetime                       datetime64[ns]
passenger_count                         int64
pickup_longitude                       float64
pickup_latitude                       float64
```

```

dropoff_longitude      float64
dropoff_latitude       float64
store_and_fwd_flag     int32
trip_duration          int64
pickup_day             int64
pickup_dayow           object
pickup_dow             int64
pickup_hour            int64
distance               float64
avg_speed_h            float64
work_hours             int32
dtype: object

```

```

In [73]: # Separating the object datatypes from the dataframe
data1 = train.select_dtypes(exclude = ['object', 'datetime'])
data1.columns

```

```

Out[73]: Index(['vendor_id', 'passenger_count', 'pickup_longitude', 'pickup_latitude',
              'dropoff_longitude', 'dropoff_latitude', 'store_and_fwd_flag',
              'trip_duration', 'pickup_day', 'pickup_dow', 'pickup_hour', 'distance',
              'avg_speed_h', 'work_hours'],
              dtype='object')

```

### Separating dependent and independent variables

```

In [74]: #seperating independent and dependent variables
train_x = data1.drop(['trip_duration', 'avg_speed_h'], axis=1)
train_y = data1['trip_duration']
train_x.shape, train_y.shape

```

```

Out[74]: ((543018, 12), (543018,))

```

```

In [75]: # Separating the object datatypes from the dataframe
data2 = test.select_dtypes(exclude = ['object', 'datetime'])
test_x = data2.drop(['trip_duration', 'median_duration', 'avg_speed_h'], axis=1)
test_y = data2['trip_duration']
test_x.shape, test_y.shape

```

```

Out[75]: ((182331, 12), (182331,))

```

```

In [76]: #importing KNN regressor and metric mse

from sklearn.neighbors import KNeighborsRegressor as KNN
from sklearn.metrics import mean_squared_error as mse

from sklearn.model_selection import ShuffleSplit
from sklearn.model_selection import cross_val_score

```

```

In [77]: # Creating instance of KNN
knn = KNN(n_neighbors = 10)

# Fitting the model
cv = ShuffleSplit(n_splits=4, test_size=0.1, random_state=10)
score = cross_val_score(knn, train_x, np.ravel(train_y), cv=cv, scoring=rmsle)
score

```

```

Out[77]: array([0.44162558, 0.44000236, 0.44092562, 0.44213926])

```

```

In [78]: score.mean()

```

Out[78]: 0.44117320433876794

In [79]: `knn.fit(train_x, train_y)`

Out[79]: `▼ KNeighborsRegressor`  
`KNeighborsRegressor(n_neighbors=10)`

In [80]: `# Predicting over the Train Set and calculating MSE`  
`test_predict = knn.predict(test_x)`  
`k = mse(np.log(test_predict+1), np.log(test_y+1), squared = False)`  
`print('Test RMSLE ', k )`

Test RMSLE 0.519649027522984

In [81]: `def Val_score(n_neighbors):`  
`'''`  
 `takes range of n_neighbors as input`  
 `returns Mean and Standard Deviation for each value of n_neighbors`  
`'''`  
 `avg = []`  
 `std = []`  
  
 `for i in n_neighbors:`  
  
 `# 10 fold cross validation for every value of n_neighbor`  
 `cv = ShuffleSplit(n_splits=4, test_size=0.1, random_state=10)`  
 `score = cross_val_score(KNN(n_neighbors = i), train_x, np.ravel(train_y), cv=cv, scor`  
 `# adding mean to avg list`  
 `avg.append(score.mean())`  
  
 `# adding standard deviation to std list`  
 `std.append(score.std())`  
  
 `return avg, std`

In [82]: `n_neighbors = range(1,20)`  
`mean, std = Val_score(n_neighbors)`


In [83]: `plt.plot(n_neighbors[1:20], mean[1:20], color = 'green', label = 'mean' )`  
`plt.xlabel('n_neighbors')`  
`plt.ylabel('Mean Score')`  
`plt.title('Mean Validation score')`

Out[83]: `Text(0.5, 1.0, 'Mean Validation score')`



```
In [84]: plt.plot(n_neighbors[1:20], std[1:20], color = 'red', label = 'Standard deviation' )
plt.xlabel('n_neighbors')
plt.ylabel('magnitude')
plt.title('Standard Deviation of Validation score')
```

```
Out[84]: Text(0.5, 1.0, 'Standard Deviation of Validation score')
```



```
In [85]: def Elbow(K):
#initiating empty list
test_rmsle = []

#training model for every value of K
for i in K:
    #Instance of KNN
    reg = KNN(n_neighbors = i)
    reg.fit(train_x, train_y)
    #Appending mse value to empty list calculated using the predictions
    tmp = reg.predict(test_x)
    tmp = mse(np.log(tmp+1),np.log(test_y+1), squared = False)
    test_rmsle.append(tmp)

return test_rmsle
```

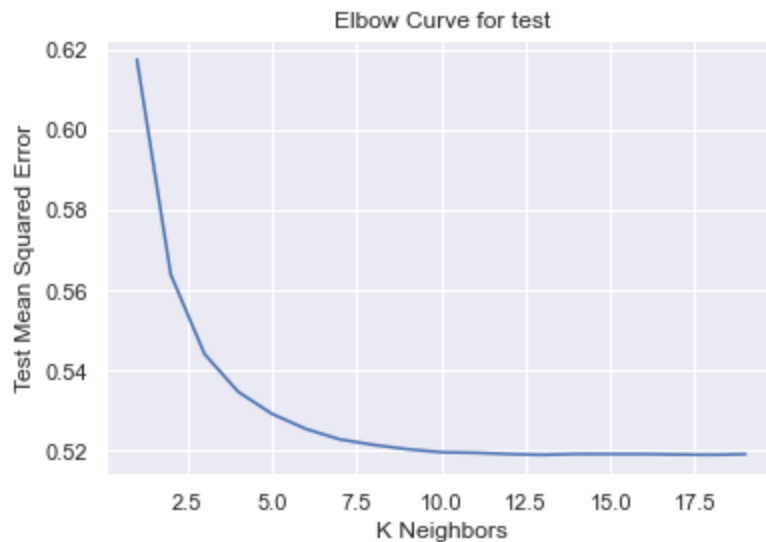
```
In [86]: #Defining K range
k = range(1,20)
# calling above defined function
```

```

b = Elbow(k)
# plotting the Curves
plt.plot(k, b)
plt.xlabel('K Neighbors')
plt.ylabel('Test Mean Squared Error')
plt.title('Elbow Curve for test')

```

Out[86]: Text(0.5, 1.0, 'Elbow Curve for test')



```

In [87]: # Creating instance of KNN with optimal value of K = 10 as it has low RMSLE for both train and test
reg = KNN(n_neighbors = 10)

# Fitting the model
cv = ShuffleSplit(n_splits=4, test_size=0.1, random_state=10)
knn_train_score = cross_val_score(reg, train_x, np.ravel(train_y), cv=cv, scoring='rmsle')
knn_train_score = knn_train_score.mean()
print('Train RMSLE      ', knn_train_score)
reg.fit(train_x, train_y)

# Predicting over the Train Set and calculating F1
test_predict = reg.predict(test_x)
knn_test_score = mse(np.log(test_predict+1), np.log(test_y+1), squared = False)
print('Test RMSLE      ', knn_test_score)

```

```

Train RMSLE      0.44117320433876794
Test RMSLE      0.519649027522984

```

## Linear Regression

```

In [90]: import sklearn
from sklearn.linear_model import LinearRegression as LR

```

```

In [91]: # Creating instance of Linear Regression
lr = LR(normalize = True)

```

```

In [92]: lr.fit(train_x, train_y)

```

Out[92]:

▼ LinearRegression  
LinearRegression(normalize=True)

```
In [93]: cv = ShuffleSplit(n_splits=4, test_size=0.1, random_state=0)
lr_score = cross_val_score(lr, train_x, np.ravel(train_y), cv=cv, scoring=rmsle)
lr_score.mean()
```

Out[93]: 0.5433030182199867

```
In [94]: # Predicting over the train set and calculating the error
train_predict = lr.predict(train_x)
train_predict[train_predict<0] = 0

lr_train_score = mse(np.log(train_predict+1), np.log(train_y+1), squared = False)
print("Training RMSLE", lr_train_score)
```

Training RMSLE 0.5395718246312002

```
In [95]: # Predicting over the Train Set and calculating F1
test_predict = lr.predict(test_x)
test_predict[test_predict<0] = 0

lr_test_score = mse(np.log(test_predict+1), np.log(test_y+1), squared = False)
print('Test RMSLE      ', lr_test_score)
```

Test RMSLE 0.6179021010187421

```
In [96]: reg_data = pd.DataFrame({'coeff': lr.coef_, 'features': train_x.columns})
reg_data
```

Out[96]:

	coeff	features
0	10.348246	vendor_id
1	4.337605	passenger_count
2	471.860923	pickup_longitude
3	-3156.478564	pickup_latitude
4	773.384686	dropoff_longitude
5	110.073430	dropoff_latitude
6	65.911018	store_and_fwd_flag
7	0.670932	pickup_day
8	12.044423	pickup_dow
9	6.992475	pickup_hour
10	99.694782	distance
11	266.542266	work_hours

**We can see from the coefficients that the latitudes, longitudes, distance and work\_hours affect the trip\_duration the most.**

```
In [97]: from sklearn.linear_model import Ridge
```



```
alpha = [0, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 25]
```

```
In [98]: # defining a function which will fit ridge regression model, plot the results, and return
def ridge_regression(train_x, train_y, test_x, test_y, alpha):
    #Fit the model
    ridgereg = Ridge(alpha=alpha, normalize=True)
    ridgereg.fit(train_x, train_y)
    train_y_pred = ridgereg.predict(train_x)
    test_y_pred = ridgereg.predict(test_x)
    train_y_pred[train_y_pred<0] = 0
    test_y_pred[test_y_pred<0] = 0

    #Return the result in pre-defined format
    k1 = mse(np.log(train_y_pred+1), np.log(train_y+1), squared = False)
    ret = [k1]

    k2 = mse(np.log(test_y_pred+1), np.log(test_y+1), squared = False)
    ret.extend([k2])

    ret.extend([ridgereg.intercept_])
    ret.extend(ridgereg.coef_)

    return ret
```

```
In [99]: ind = ['alpha_%.2g'%alpha[i] for i in range(0,10)]
rmsle_matrix_ridge = pd.DataFrame(index = ind, columns = ['train_rmsle', 'test_rmsle', 'in
```

```
In [100... for i in range(10):
    rmsle_matrix_ridge.iloc[i,] = ridge_regression(train_x, train_y, test_x, test_y, alpha
```

```
In [101... #Set the display format to be scientific for ease of analysis
pd.options.display.float_format = '{:,.5g}'.format
rmsle_matrix_ridge
```

```
Out[101...
```

	train_rmsle	test_rmsle	intercept	coef_Var_1	coef_Var_2	coef_Var_3	coef_Var_4	coef_Var_5	coef_V
alpha_0	0.53957	0.6179	2.165e+05	10.348	4.3376	471.86	-3,156.5	773.38	11
alpha_1e-08	0.53957	0.6179	2.165e+05	10.348	4.3376	471.86	-3,156.5	773.38	11
alpha_0.0001	0.53959	0.61792	2.1655e+05	10.347	4.3376	472.44	-3,156	773.44	10
alpha_0.001	0.53975	0.61809	2.1699e+05	10.339	4.3375	477.67	-3,151.5	773.9	10
alpha_0.01	0.54138	0.6198	2.2124e+05	10.256	4.3358	528.41	-3,107.4	778.45	56
alpha_1	0.63811	0.71384	2.7084e+05	6.3815	3.0161	1,481.5	-1,538.5	840.06	-87
alpha_5	0.71733	0.78836	1.3127e+05	2.6802	1.1992	754.56	-597.8	434.3	-44
alpha_10	0.74064	0.81027	77,840	1.5507	0.68291	447.85	-341.2	260.75	-2
alpha_20	0.75583	0.82452	43,032	0.8411	0.36698	246.01	-183.69	144.27	-14
alpha_25	0.75924	0.82775	35,246	0.68445	0.29805	200.7	-149.25	117.88	-11

```
In [102... #Importing Lasso model from sklearn's linear_model module
from sklearn.linear_model import Lasso
```

```
In [103... alpha_lasso = [0, 1e-10, 1e-8, 1e-5, 1e-4, 1e-3, 1e-2, 1, 5, 10]
```

```
In [104... # defining a function which will fit lasso regression model, plot the results, and return
def lasso_regression(train_x, train_y, test_x, test_y, alpha):
    #Fit the model
    lassoreg = Lasso(alpha=alpha, normalize=True)
    lassoreg.fit(train_x, train_y)
    train_y_pred = lassoreg.predict(train_x)
    test_y_pred = lassoreg.predict(test_x)
    train_y_pred[train_y_pred<0] = 0
    test_y_pred[test_y_pred<0] = 0

    #Return the result in pre-defined format
    k1 = mse(np.log(train_y_pred+1), np.log(train_y+1), squared = False)
    ret = [k1]

    k2 = mse(np.log(test_y_pred+1), np.log(test_y+1), squared = False)
    ret.extend([k2])

    ret.extend([lassoreg.intercept_])
    ret.extend(lassoreg.coef_)

    return ret
```

```
In [105... rmsle_matrix_lasso = pd.DataFrame(index = ind, columns = ['train_rmsle', 'test_rmsle', 'ir
```

```
In [106... for i in range(10):
    rmsle_matrix_lasso.iloc[i,] = lasso_regression(train_x, train_y, test_x, test_y, alpha
```

```
In [107... rmsle_matrix_lasso
```

	train_rmsle	test_rmsle	intercept	coef_Var_1	coef_Var_2	coef_Var_3	coef_Var_4	coef_Var_5	coef_V
alpha_0	0.53957	0.6179	2.165e+05	10.348	4.3376	471.86	-3,156.5	773.38	11
alpha_1e-08	0.53957	0.6179	2.165e+05	10.348	4.3376	471.86	-3,156.5	773.38	11
alpha_0.0001	0.53957	0.6179	2.165e+05	10.348	4.3376	471.86	-3,156.5	773.38	11
alpha_0.001	0.53957	0.6179	2.1649e+05	10.336	4.3335	471.86	-3,156.1	773.27	10
alpha_0.01	0.53959	0.61792	2.1641e+05	10.223	4.2963	471.88	-3,152.3	772.22	10
alpha_1	0.53974	0.6181	2.156e+05	9.098	3.9247	472.03	-3,115.1	761.76	70
alpha_5	0.5421	0.62043	1.9629e+05	0	0.023554	437.87	-2,875.1	627.27	
alpha_10	0.77474	0.8422	847.08	0	0	0	-0	0	
alpha_20	0.77474	0.8422	847.08	0	0	0	-0	0	
alpha_25	0.77474	0.8422	847.08	0	0	0	-0	0	

## Decision Tree

```
In [108... from sklearn.tree import DecisionTreeRegressor as DT
```

```
dreg = DT()
```

```
In [109... dreg = dreg.fit(train_x,np.ravel(train_y))
```

```
In [110... train_predict = dreg.predict(train_x)

dt_train_score = mse(np.log(train_predict+1), np.log(train_y+1), squared = False)
print('Train RMSLE      ', dt_train_score)
```

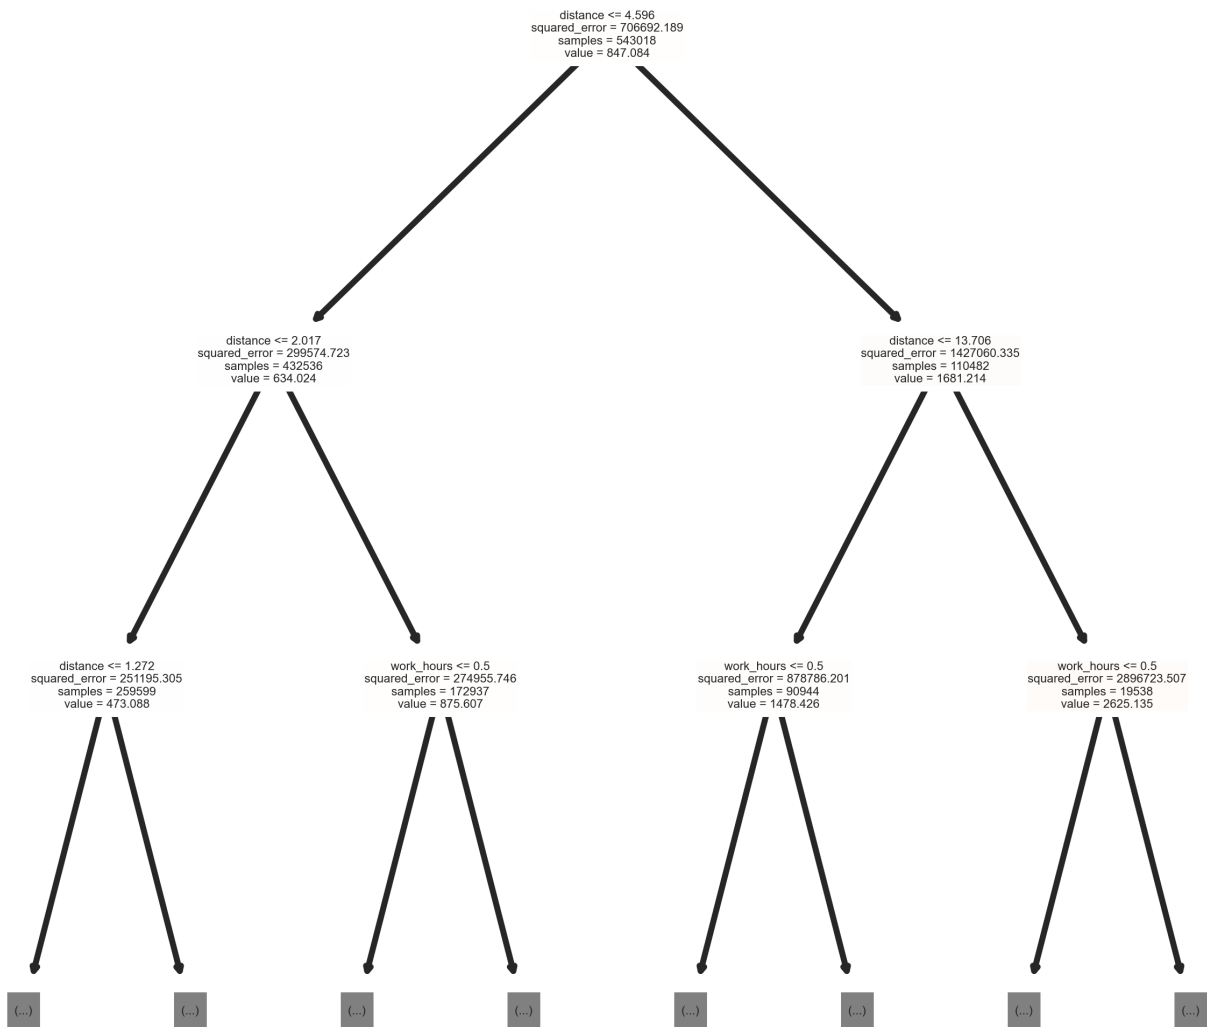
```
Train RMSLE      0.0015605166604175274
```

```
In [111... test_predict = dreg.predict(test_x)

dt_test_score = mse(np.log(test_predict+1), np.log(test_y+1), squared = False)
print('Test RMSLE      ',dt_test_score)
```

```
Test RMSLE      0.568099778720311
```

```
In [112... from sklearn import tree
fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (4,4), dpi=800)
tree.plot_tree(dreg,
                feature_names = train_x.columns,
                max_depth = 2,
                filled = True);
```



- We can see from the decision tree that the distance and work\_hours are the two most important features to predict the trip\_duration.

---

## Random Forest

---

In [116...

```
from sklearn.ensemble import RandomForestRegressor as RF
rf = RF()
cv = ShuffleSplit(n_splits=4, test_size=0.1, random_state=0)
rf_score = cross_val_score(rf, train_x, np.ravel(train_y), cv=cv, scoring=rmsle)
rf_score.mean()
```

Out[116...

0.36326167406351806

In [117...

```
rf = rf.fit(train_x, np.ravel(train_y))
```

In [118...

```
train_predict = rf.predict(train_x)

rf_train_score = mse(np.log(train_predict+1), np.log(train_y+1), squared = False)
print('Train RMSLE      ', rf_train_score)
```

Train RMSLE 0.15588947391892835

In [119...

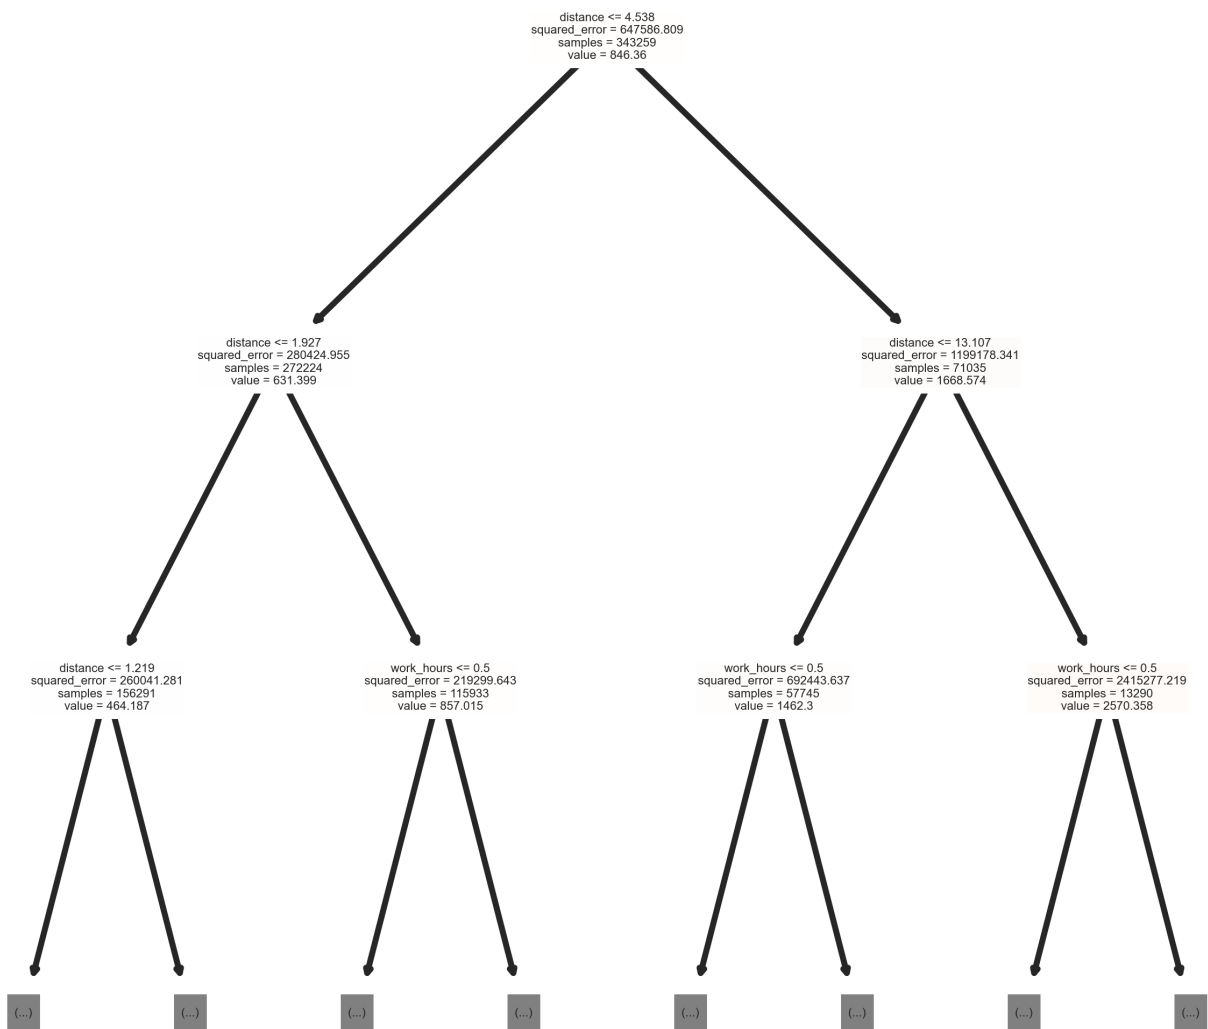
```
test_predict = rf.predict(test_x)

rf_test_score = mse(np.log(test_predict+1), np.log(test_y+1), squared = False)
print('Test RMSLE      ', rf_test_score)
```

Test RMSLE 0.4803150083992271

In [120...

```
fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (4,4), dpi=800)
tree.plot_tree(rf.estimators_[0],
               feature_names = train_x.columns,
               max_depth = 2,
               filled = True);
```



```

In [121...
result = pd.DataFrame(data = {'Model' : ['KNN', 'Linear Reg', 'Decision Tree', 'Random Forest'],
                             'train_RMSLE': [knn_train_score, lr_train_score, dt_train_score, rf_train_score],
                             'test_RMSLE': [knn_test_score, lr_test_score, dt_test_score, rf_test_score]})
result

```

Out[121...

	Model	train_RMSLE	test_RMSLE
0	KNN	0.44117	0.51965
1	Linear Reg	0.53957	0.6179
2	Decision Tree	0.0015605	0.5681
3	Random Forest	0.15589	0.48032

```

In [122...
tidy = result.melt(id_vars='Model').rename(columns={str.title: str.lower})
tidy

```

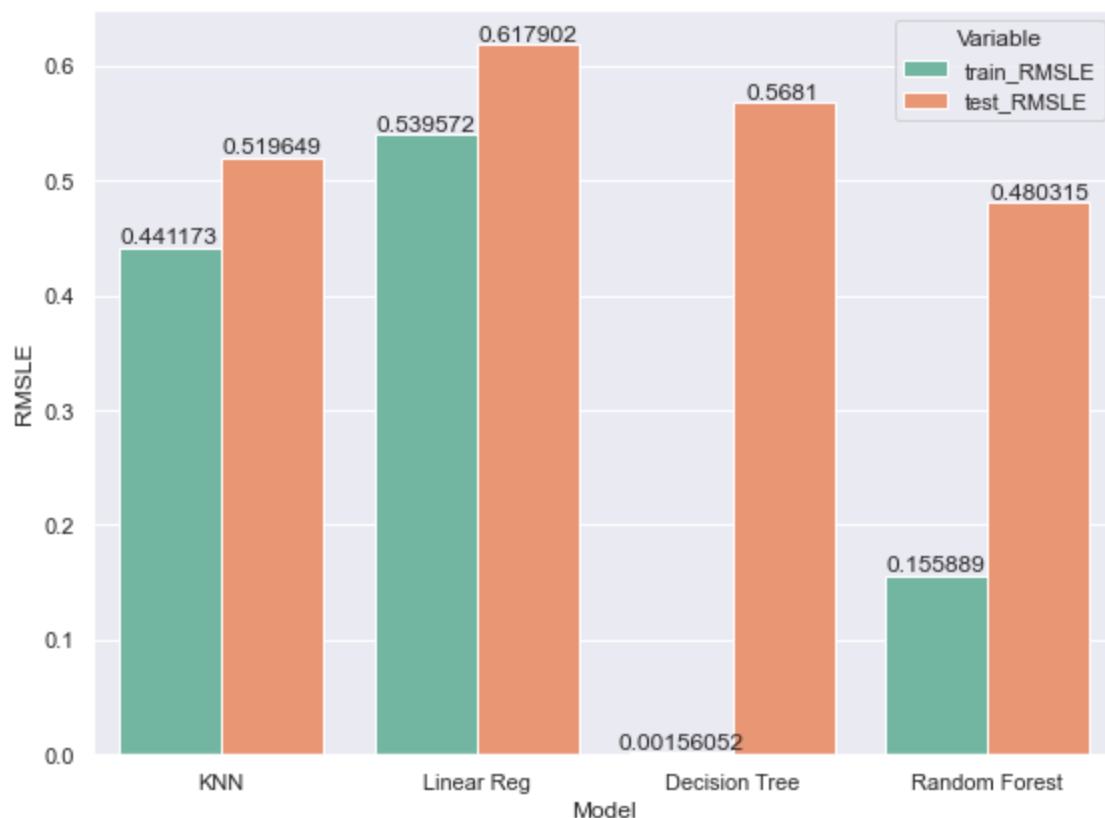
Out[122...

	Model	Variable	Value
0	KNN	train_RMSLE	0.44117
1	Linear Reg	train_RMSLE	0.53957
2	Decision Tree	train_RMSLE	0.0015605
3	Random Forest	train_RMSLE	0.15589
4	KNN	test_RMSLE	0.51965
5	Linear Reg	test_RMSLE	0.6179
6	Decision Tree	test_RMSLE	0.5681
7	Random Forest	test_RMSLE	0.48032

In [123...

```
fig,ax = plt.subplots(figsize = (8,6))
sns.barplot(x = 'Model', y = 'Value', hue = 'Variable', data = tidy, palette = 'Set2')
ax.bar_label(container=ax.containers[0])
ax.bar_label(container=ax.containers[1])
ax.set_ylabel('RMSLE')

fig.tight_layout()
```



- We can see from the results that decision tree model gives the lowest RMSLE for our training data among other models but a higher RMSLE than KNN and Random Forest for the test data. This is because the models overfits to the training data and hence does not perform very well on the test data.
- The best model among the four would be the Random Forest model as it gives very low RMSLE for training data and the lowest RMSLE for the test data.